

Accelerating Low-Frequency Convergence for Limited-Angle DBT

Two-Channel Fidelity in PDHG

Taro Iyadomi¹ Ricardo Parada² Anna Kim³ Lily Jiang⁴ Emil Sidky⁵ William Chang⁴

CT Meeting 2026

¹Northwestern University ²UC Riverside ³UC San Diego

⁴UC Los Angeles ⁵University of Chicago

Motivation

What is limited-angle reconstruction?

- Tomography recovers a cross-section from X-ray projections taken at many source angles around the object.
- **Sufficient sampling.** For a fan-beam scan, a source arc of 180° plus the fan angle already measures every line integral — a full 360° rotation only re-measures each ray from the opposite side.
- **Limited-angle reconstruction (LAR)** uses a source arc *below* that threshold. The measurements no longer determine the image: a wedge of spatial frequencies is never observed.
- Digital breast tomosynthesis (DBT) operates deep in this regime — a source arc of only $[15^\circ, 50^\circ]$ — so its reconstruction is severely ill-posed.

Why limited-angle?

- **Dose** — fewer views means less ionizing radiation per scan.
- **Geometry** — in DBT the detector stays under the compressed breast while the X-ray tube sweeps a short arc; a full gantry rotation is mechanically impossible.
- **Acquisition time** — a short sweep is fast, limiting motion blur and patient discomfort.
- **Use cases** — DBT for early breast-cancer detection; also lung-nodule and kidney-stone imaging.

*[DBT geometry schematic
50° source arc,
detector, breast cross-section]*

Two reconstruction settings: 2D and 3D

The two-channel method is evaluated in both geometries:

- **2D fan-beam** — a single slice; cheap enough to sweep image resolution and noise level.
- **3D cone-beam** — the actual DBT geometry; full-volume reconstruction on anatomical phantoms.

Both are severely under-determined: the number of unknowns far exceeds the number of measurements, increasingly so at finer resolution.

Slice resolution	Unknowns (pixels)	Measurements (rays)
512^2	262,144	25,600
256^2	65,536	25,600
128^2	16,384	25,600

Measurements = 25 views $\times$ 1024 detector bins.

The reconstruction setup

Reconstruction with PDHG (Chambolle–Pock)

The image is found by minimizing directional-TV and ℓ_1 sparsity, subject to the projections matching the data within a tolerance, with $f \geq 0$.

- **Why not gradient descent?** The objective is non-differentiable (ℓ_1 , TV) and hard-constrained — gradient descent needs a smooth, unconstrained objective.
- **PDHG** handles this class with cheap projections and proximal maps, no derivatives. It keeps a primal variable (the image) and a dual variable (the data residual), and alternates:
 - a **dual step** — pull the data residual toward the tolerance ball (step σ);
 - a **primal step** — pull the image toward sparsity and $f \geq 0$ (step τ).
- τ and σ are coupled through the operator norm $\|K\|$; we set them by an empirical balance heuristic (appendix).

Single-channel data fidelity

Prior work¹ solves a constrained DTV / ℓ_1 problem with a square-root Hanning filter on the sinogram residual:


$$\begin{aligned} \min_{f \geq 0} \quad & \alpha_x \|\partial_x f\|_1 + \alpha_y \|\partial_y f\|_1 + \alpha_z \|\partial_z f\|_1 + \beta \|f\|_1 \\ \text{s.t.} \quad & \|R_{[c]}(g - Xf)\|_2 \leq \varepsilon \sqrt{|g|}. \end{aligned}$$

- $R_{[c]} = \text{Hann}^{1/2}(c)$ attenuates low spatial frequencies in the residual — noise suppression by design.
- Solved with PDHG and He–Yuan over-relaxation ($\rho \approx 1.75$).

¹E. Y. Sidky et al., *J. Med. Imaging* 12(S1):S13013, 2025.

The problem

What happens during a single-channel run



- Fine structure (calcifications, glandular speckling) settles fast.
- Large soft-tissue regions stay blurred, with a slowly-varying intensity error, through iter ~ 100 – 300 .
- This is the low-frequency convergence bottleneck.

2D fan-beam · 25 views · 50° arc · 256² grid

Why does this happen?

In the detector-frequency domain, the data-fidelity weight is

$$W(g - Xf), \quad W = R_{[c]} = \text{Hann}^{1/2}(c).$$

- PDHG corrects mode k at a rate $\propto s_k^2$ (squared singular values of WX).
- In LAR DBT, s_k is already tiny for low k (ill-conditioning).
- $\text{Hann}^{1/2}$ shrinks low- k further *by design*.
- Stability bound is dictated by the *largest* s_k (high frequencies).

Consequence

The optimizer is forced to use step sizes that under-drive the very modes that are hardest to recover.

The idea: two-channel fidelity


Split the sinogram into two bands

$$\begin{aligned} & \min_{f \geq 0} \alpha_{\bullet} \|\partial_{\bullet} f\|_1 + \beta \|f\|_1 \\ \text{s.t.} \quad & \begin{cases} \|R_{\text{hi}}(g - Xf)\|_2 \leq \varepsilon_{\text{hi}} \sqrt{|g|}, \\ \|R_{\text{lo}}(g - Xf)\|_2 \leq \varepsilon_{\text{lo}} \sqrt{|g|}. \end{cases} \end{aligned}$$

- $R_{\text{hi}} = \text{Hann}^{1/2}(c)$ – preserves the high-frequency channel from prior work.
- $R_{\text{lo}} = \text{Hann}^{1/2}(c_{\text{lo}})$ with $c_{\text{lo}} \gg c$ – a narrow-band low-pass channel that targets the slow modes.
- Independent dual variables $y_{\text{hi}}, y_{\text{lo}}$ with independent step sizes $\sigma_{\text{hi}}, \sigma_{\text{lo}}$.

Step sizes follow the same balance heuristic as the single-channel paper: $\tau \sigma_{\text{hi}} \|K\|^2 \approx 1$, with $\sigma_{\text{lo}} = 4\sigma_{\text{hi}}$ to drive the slow modes harder.


The two-channel iteration



- The data residual is split into a high- and a low-frequency band.
- Each band carries its own dual variable and its own step size.
- The low band is driven $4\times$ harder ($\sigma_{lo} = 4 \sigma_{hi}$) — the slow modes receive the larger correction.
- The primal update and the image itself are unchanged from single-channel.

Results


Iteration ladder — single vs. two-channel



- 256×256 , 25 views over 50° .
- Two-channel resolves recognizable soft-tissue structure several iterations earlier.

2D fan-beam · 25 views · 50° arc · 256^2 grid


Soft-tissue ROI: low-frequency error



Iteration ladder cropped to a soft-tissue ROI. Single-channel carries a slowly-varying intensity error through iterations 50–200; two-channel reaches a uniform soft-tissue level several iterations earlier.


$$2D \text{ fan-beam} \cdot 25 \text{ views} \cdot 50^\circ \text{ arc} \cdot 256^2 \text{ grid}$$

Where is the error?




Absolute error $|\text{reconstruction} - \text{ground truth}|$ (scale bar below; darker is less error). Single-channel at iter 100–200 shows a broad bright low-frequency error region; the two-channel row stays darker throughout.

Convergence behaviour




- Single-channel oscillates strongly through iter ~ 200 — the cross-band interference predicted by the spectral analysis.
- Two-channel converges smoothly from the start.

Persistence across image sizes




- Two-channel sits below single-channel across all three grids.
- Iter-500 RMSE reduction: 19% at 512^2 , 22% at 256^2 , 62% at 128^2 .
- Consistent across the resolution sweep.

Design knob 1: band selection (2D)




Cutoff sweep at fixed $c_{hi} = 4$, c_{lo} varying from 4 (PoU) to 32. The two-channel image-RMSE curves are visually indistinguishable across the sweep — the method is robust to filter shape, and PoU vs non-PoU is empirically equivalent.

Design knob 2: LF tolerance (2D) — a clear sweet spot




Per-band tolerance sweep, $r = \varepsilon_{lo}/\varepsilon_{hi}$, at $c_{hi} = 4$, $c_{lo} = 8$. A broad interior optimum around $r \approx 15$; the paper's $r = 1.25$ is too tight; $r \rightarrow \infty$ collapses (no LF constraint — the recon flat-lines near the single-channel value).

Design knob 2: tuned $r = 15$ recovers earlier (2D)



Single-channel vs paper config ($r = 1.25$) vs tuned ($r = 15$). At iter 100 the tuned RMSE is 0.122 vs paper 0.162 (-25%) and vs single 0.220 (-45%). Loosening the LF constraint lets the low-frequency channel correct the bias.


3D: Analytic breast phantom



Analytic breast phantom: fatty soft-tissue background with discrete fatty-mass inserts and a chest-wall region. $144 \times 144 \times 32$ grid at 1.5 mm isotropic voxels.

3D cone-beam · 25 views · 50° arc · $144 \times 144 \times 32$ grid


Analytic breast — mid-axial iteration ladder



Cone-beam LAR (25 views / 50°), mid-axial slice. Two-channel image RMSE is 8% lower at iter 20, rising to a 35% reduction at iter 200; it resolves the discrete masses earlier.


3D cone-beam · 25 views · 50° arc · $144 \times 144 \times 32$ grid

Analytic breast — soft-tissue ROI



Lateral soft-tissue ROI with several fatty-mass inserts. Image RMSE at iter 100: single 0.077 vs. two-channel 0.051 — a 34% reduction. Two-channel shows discrete mass contours while single-channel retains a diffuse intensity bias.

Analytic breast — reconstruction error



Absolute error $|\text{reconstruction} - \text{ground truth}|$, full mid-axial slice (scale bar below; darker is less error).
The two-channel row stays consistently darker than single-channel.

3D cone-beam · 25 views · 50° arc · 144 × 144 × 32 grid

Analytic breast — convergence





Image RMSE vs. iteration. Two-channel leads throughout: a 34% reduction at iter 100, peaking near 35% at iter 200 and holding 16% at iter 500.


Design knob 1: band selection (3D breast)



Same cutoff sweep on the 3D analytic breast at 50° arc, c_{lo} in $\{4, 6, 8, 12, 16\}$. The five curves overlap throughout — the band-selection robustness carries over from 2D.


3D cone-beam · 25 views · 50° arc · $144 \times 144 \times 32$ grid

Design knob 2: LF tolerance (3D breast) — paper config near-optimal



r sweep on the 3D analytic breast. Unlike the 2D case, the paper's $r = 1.25$ is already at or near the iter-100⁺ optimum — loosening to $r = 15$ helps only at iter ~ 50 and degrades thereafter. The 3D semi-convergence rise dominates the picture.


Wider arcs: head at 100°



Analytic head, full mid-axial slice. At 100° the streaking recedes and the two-channel improvement is visible directly: the two-channel row is markedly cleaner than single-channel.

3D cone-beam · 25 views · 100° arc · $176 \times 176 \times 16$ grid


Wider arcs: jaw at 100°



Analytic jaw, full mid-axial slice. Same picture as head: at 100° the two-channel row is visibly cleaner than single-channel.

3D cone-beam · 25 views · 100° arc · $144 \times 144 \times 80$ grid

Convergence across arcs — head and jaw



Single-channel (dashed) vs two-channel (solid) at 50° (red), 75° (green), 100° (blue). At every arc the two-channel curve sits below single-channel — the gap is most legible at 100° (where the visual ladder above confirms it) and persists at 75° and 50°, where the visual difference is subtler.

3D cone-beam · 25 views · analytic head and jaw

Conclusion


Summary

- LAR DBT reconstruction has a *mode-dependent* convergence bottleneck: low spatial frequencies are starved of corrective drive because the PDHG step size is bound by high-frequency stability.
- Splitting the data fidelity into a low-pass and a high-pass channel, each with its own dual variable and step size, accelerates low-frequency settling without breaking stability.
- Confirmed on the 2D paper breast and the 3D analytic breast at a clinical 50° arc.
- **Design space.** Band selection is robust on both phantoms (filter shape doesn't matter much). LF tolerance is the real lever — tunable on 2D (best $r \approx 15$ gives a 25% iter-100 reduction over the paper config and 45% over single-channel), near-optimal at the paper's $r = 1.25$ on 3D breast.

- **Multi-band fidelity.** Generalize from two channels to a geometric sequence of band-pass filters with matched per-shell step sizes. Manuscript in preparation.
- **Theorem-compliant step sizes.** Derive step sizes from a weighted power iteration so the empirical balance heuristic connects to the product-space PDHG convergence guarantee.
- **LF-tolerance scaling.** The 2D paper-breast optimum $r \approx 15$ does not transfer to 3D analytic breast; characterising which problem features control the optimum is open.
- **Realistic voxelized anatomy.** On the VICTRE breast phantom the two-channel gain at 50° is front-loaded — present early, then overtaken; sustaining it is ongoing.

Questions?

Backup: 2D ϵ_{hi} scale sweep




With the LF-tolerance ratio held at $r = 10$, the overall scale ϵ_{hi}/ϵ has an interior optimum near 1–2. Tightening to 0.5 hurts early iterations; loosening to 4 degrades uniformly. The paper's $\epsilon_{hi} = \epsilon$ is near-optimal.

2D fan-beam · 25 views · 50° arc · 256² grid

Backup: 3D analytic head and jaw


3D: Analytic head phantom



Analytic head phantom: skull shell with sinuses, brain with ventricle and FORBILD low-contrast spots, subdural hematoma. 26 cm FOV, thin axial slab.

3D cone-beam · 25 views · 50° arc · $176 \times 176 \times 16$ grid


Analytic head — mid-axial iteration ladder



Full mid-axial slice, single vs two-channel. Both methods drift upward as iterations proceed under the under-determined geometry; the magnitude difference image and convergence plot make the comparison quantitative.

3D cone-beam · 25 views · 50° arc · 176 × 176 × 16 grid

Analytic head — reconstruction error



Absolute error $|\text{reconstruction} - \text{ground truth}|$, full mid-axial slice (scale bar below; darker is less error).
The two-channel row stays consistently darker — less error — than single-channel.

3D cone-beam · 25 views · 50° arc · 176 × 176 × 16 grid

Analytic head — convergence





Image RMSE vs. iteration. Two-channel stays below single-channel from early iterations onward; the gap widens from 12% at iter 100 to 19% at iter 500.

3D cone-beam · 25 views · 50° arc · $176 \times 176 \times 16$ grid


3D: Analytic jaw phantom



Analytic jaw phantom: mandible/maxilla with 28 teeth, three gold crowns (high- μ scatterers), low-contrast tongue tumour. 12 cm tall volume.

3D cone-beam · 25 views · 50° arc · 144 × 144 × 80 grid


Analytic jaw — mid-axial iteration ladder



Full mid-axial slice, single vs two-channel. The gold crowns ($\mu = 19.3$) dominate the projection energy; the magnitude difference image and convergence plot quantify the gap.

3D cone-beam · 25 views · 50° arc · 144 × 144 × 80 grid

Analytic jaw — reconstruction error



Absolute error $|\text{reconstruction} - \text{ground truth}|$, full mid-axial slice (scale bar below; darker is less error).
The two-channel row stays darker — less error — than single-channel.

3D cone-beam · 25 views · 50° arc · 144 × 144 × 80 grid

Analytic jaw — convergence





Image RMSE vs. iteration. Two-channel converges to a lower RMSE; the gap grows from 5% at iter 100 to 17% at iter 500.

3D cone-beam · 25 views · 50° arc · $144 \times 144 \times 80$ grid

Backup: wider-arc reconstructions


Analytic breast at 75° — iteration ladder



Full mid-axial slice, 75° arc. Two-channel resolves the discrete masses earlier; iter-500 RMSE is 16% lower (single 0.069, two-channel 0.058).

3D cone-beam · 25 views · 75° arc · $144 \times 144 \times 32$ grid

Analytic breast at 75° — reconstruction error



Absolute error $|\text{reconstruction} - \text{ground truth}|$, 75° arc (darker is less error). The two-channel row stays darker than single-channel.

3D cone-beam · 25 views · 75° arc · $144 \times 144 \times 32$ grid

Analytic breast at 75° — convergence





Image RMSE vs. iteration, 75° arc. Two-channel stays below single-channel; the gap is +30% at iter 200 and +16% at iter 500.

3D cone-beam · 25 views · 75° arc · $144 \times 144 \times 32$ grid


Analytic head at 75° — iteration ladder



Full mid-axial slice, 75° arc — visibly cleaner than at 50° . Two-channel iter-500 RMSE is 17% lower (single 0.649, two-channel 0.541).

3D cone-beam · 25 views · 75° arc · $176 \times 176 \times 16$ grid

Analytic head at 75° — reconstruction error



Absolute error $|\text{reconstruction} - \text{ground truth}|$, 75° arc (darker is less error). The two-channel row is consistently darker.

3D cone-beam · 25 views · 75° arc · $176 \times 176 \times 16$ grid

Analytic head at 75° — convergence





Image RMSE vs. iteration, 75° arc. Two-channel stays below single-channel (+19% at iter 200, +17% at iter 500); the semi-convergence rise is smaller than at 50°.

3D cone-beam · 25 views · 75° arc · $176 \times 176 \times 16$ grid


Analytic jaw at 75° — iteration ladder



Full mid-axial slice, 75° arc. Two-channel iter-500 RMSE is 14% lower (single 0.593, two-channel 0.509).

3D cone-beam · 25 views · 75° arc · $144 \times 144 \times 80$ grid

Analytic jaw at 75° — reconstruction error



Absolute error $|\text{reconstruction} - \text{ground truth}|$, 75° arc (darker is less error). The two-channel row stays darker.

3D cone-beam · 25 views · 75° arc · $144 \times 144 \times 80$ grid

Analytic jaw at 75° — convergence





Image RMSE vs. iteration, 75° arc. Two-channel stays below single-channel; the gap reaches +14% at iter 500.

3D cone-beam · 25 views · 75° arc · $144 \times 144 \times 80$ grid


Analytic breast at 100° — iteration ladder



Full mid-axial slice, 100° arc. At this wide arc the breast problem is well-conditioned — both methods reconstruct it well, and single-channel converges fastest here (see convergence).

3D cone-beam · 25 views · 100° arc · $144 \times 144 \times 32$ grid

Analytic breast at 100° — reconstruction error



Absolute error $|\text{reconstruction} - \text{ground truth}|$, 100° arc (darker is less error). Both rows are dark — low error — with single-channel the cleaner of the two at this well-conditioned arc.

3D cone-beam · 25 views · 100° arc · $144 \times 144 \times 32$ grid


Analytic breast at 100° — convergence



Image RMSE vs. iteration, 100° arc. The breast problem is now well-conditioned: single-channel converges to a near-perfect RMSE (0.010), beating two-channel (0.021). The two-channel gain is a limited-angle effect — it fades once the arc is wide enough.

3D cone-beam · 25 views · 100° arc · $144 \times 144 \times 32$ grid


Analytic head at 100° — iteration ladder



Full mid-axial slice, 100° arc — markedly cleaner than at 50° . Two-channel iter-500 RMSE is 13% lower (single 0.570, two-channel 0.497).

3D cone-beam · 25 views · 100° arc · $176 \times 176 \times 16$ grid

Analytic head at 100° — reconstruction error



Absolute error $|\text{reconstruction} - \text{ground truth}|$, 100° arc (darker is less error). The two-channel row stays darker.

3D cone-beam · 25 views · 100° arc · $176 \times 176 \times 16$ grid

Analytic head at 100° — convergence





Image RMSE vs. iteration, 100° arc. Two-channel stays below single-channel throughout (+16% at iter 200, +13% at iter 500).

3D cone-beam · 25 views · 100° arc · $176 \times 176 \times 16$ grid


Analytic jaw at 100° — iteration ladder



Full mid-axial slice, 100° arc. Two-channel iter-500 RMSE is 8% lower (single 0.516, two-channel 0.476).

3D cone-beam · 25 views · 100° arc · $144 \times 144 \times 80$ grid

Analytic jaw at 100° — reconstruction error



Absolute error $|\text{reconstruction} - \text{ground truth}|$, 100° arc (darker is less error). The two-channel row stays darker.

3D cone-beam · 25 views · 100° arc · $144 \times 144 \times 80$ grid

Analytic jaw at 100° — convergence




Image RMSE vs. iteration, 100° arc. Two-channel stays below single-channel; the gap narrows to +8% at iter 500 as the problem becomes better conditioned.

3D cone-beam · 25 views · 100° arc · $144 \times 144 \times 80$ grid

Backup: Shepp-Logan and VICTRE

Backup: 3D phantoms at a clinical 50° arc

Two-channel vs. single-channel image-RMSE reduction (%) on the two realistic 3D phantoms, 25 views over 50°. Positive favours two-channel.


Phantom	Step-size config	iter 100	iter 200	iter 300	iter 400
Shepp-Logan	paper config	+0.3	-5.2	-6.5	-1.0
	tuned (best)	-1.5	-3.2	+0.1	-1.6
VICTRE breast	paper config	-5.6	-11.2	-17.8	-22.4
	tuned (best)	+3.3	+1.1	-6.0	-12.5

- At a realistic 50° arc neither 3D phantom shows the persistent gain seen on the analytic phantoms.
- VICTRE: tuning the CP step sizes yields a front-loaded gain — faster early, overtaken after iter ~200. Shepp-Logan is a tie: a well-conditioned negative control.
- “Tuned” reduces the 3D step-size handicap (σ -inflation 2.0 $\rightarrow$ 1.5); the gain grows as the arc narrows below 50°.

Backup: 3D Shepp-Logan phantom

Shepp-Logan phantom (10 ellipsoids, ~ 10 cm half-extent)


axial (xy)



coronal (xz)



sagittal (yz)



Classical 3D Shepp-Logan: 10 ellipsoids, skull at $\mu = 2.0$, brain at $\mu = 1.0$, three small low-contrast inserts at $\mu = 1.0 \pm 0.02$.

3D cone-beam · 25 views · 50° arc · 128 × 128 × 64 grid

Backup: Shepp-Logan at a clinical 50° arc





Image RMSE vs. iteration. Single- and two-channel are essentially superimposed (single is $\sim 7\%$ lower at iter 500). With little intermediate spatial structure there is no low-frequency deficit for the spectral split to address — a well-conditioned negative control.

3D cone-beam · 25 views · 50° arc · $128 \times 128 \times 64$ grid


Backup: 3D VICTRE breast phantom

VICTRE voxelized breast phantom (0.4 mm after 4x downsampling)


axial (xy)



coronal (xz)




sagittal (yz)



Voxelized VICTRE breast phantom (Reiser–Nishikawa anatomic model, $\sim 10^6$ voxels at 0.4 mm after 4x downsampling). Anatomically distributed glandular structure rather than discrete analytic inserts.

3D cone-beam · 25 views · 50° arc · VICTRE breast, 4x downsampled

Backup: VICTRE at a clinical 50° arc



Iteration ladder restricted to a glandular ROI; severe LAR streaking dominates both methods. Under the paper step sizes single-channel has the lower RMSE; tuning the CP step sizes produces an early two-channel lead that fades after iter ~ 200 (see summary table).

3D cone-beam · 25 views · 50° arc · VICTRE breast, $4\times$ downsampled

Backup: method details

Backup: full constrained objective

$$\begin{aligned} \min_{f \geq 0} \quad & \alpha_x \|\partial_x f\|_1 + \alpha_y \|\partial_y f\|_1 + \alpha_z \|\partial_z f\|_1 + \beta \|f\|_1 \\ \text{s.t.} \quad & \begin{cases} \|R_{\text{hi}}(g - Xf)\|_2 \leq \varepsilon_{\text{hi}} \sqrt{|g|}, \\ \|R_{\text{lo}}(g - Xf)\|_2 \leq \varepsilon_{\text{lo}} \sqrt{|g|}. \end{cases} \end{aligned}$$

Parameter choices used for the figures in this talk:

	128 ²	256 ² , 512 ²
α (DTV)	1.95	1.9, 1.7
β (ℓ_1)	10	10, 5
c_{hi}	4	4
c_{lo}	8	8
$\sigma_{\text{lo}}/\sigma_{\text{hi}}$	4	4
$\varepsilon_{\text{lo}}/\varepsilon_{\text{hi}}$	1.25	1.25

Backup: Algorithm 1 (PDHG with split fidelity)

Inputs: $g, X, R_{\text{hi}}, R_{\text{lo}}, \alpha_{x,y,z}, \beta, \varepsilon_{\text{hi,lo}}, \sigma_{\text{hi,lo}}, \tau, \rho$.

Initialize $f^0, \bar{f}^0 = f^0, y_{\text{hi}}^0 = y_{\text{lo}}^0 = 0$.

for $k = 0, 1, 2, \dots$ **do**

$$y_{\text{hi}}^{k+1} \leftarrow \text{Proj}_{B_{\text{hi}}}(y_{\text{hi}}^k + \sigma_{\text{hi}} R_{\text{hi}}(g - X \bar{f}^k)),$$

$$y_{\text{lo}}^{k+1} \leftarrow \text{Proj}_{B_{\text{lo}}}(y_{\text{lo}}^k + \sigma_{\text{lo}} R_{\text{lo}}(g - X \bar{f}^k)),$$

$$f^{k+1} \leftarrow \text{prox}_{\tau\Phi}(f^k - \tau X^\top (R_{\text{hi}}^\top y_{\text{hi}}^{k+1} + R_{\text{lo}}^\top y_{\text{lo}}^{k+1})),$$

$$\bar{f}^{k+1} \leftarrow f^{k+1} + \rho(f^{k+1} - f^k).$$

Here Φ collects the DTV $+\ell_1$ + non-negativity terms.

Backup: Step-size heuristic (and future work)

These results use the paper's empirical balance heuristic:

$$\sigma_{\text{hi}} = c_{\text{bal}}/\|K\|, \quad \tau = 1/(c_{\text{bal}}\|K\|), \quad \sigma_{\text{lo}} = 4\sigma_{\text{hi}},$$

with $\|K\|$ from an *unweighted* joint power iteration over $[R_{\text{hi}}X; R_{\text{lo}}X; \partial_x; \partial_y; \partial_z; \ell_1 I]$.

The product-space CP theorem demands $\tau \cdot \|\Sigma^{1/2}K\|^2 < 1$, which the heuristic does not strictly satisfy when $\sigma_{\text{lo}} > \sigma_{\text{hi}}$.

- Empirically stable for 2D at 500 iters; in 3D we scale $\|K\|$ up by $\sqrt{\sigma_{\text{lo}}/\sigma_{\text{hi}}}$ for stability without changing the LF/HF ratio.
- He–Yuan over-relaxation $\rho \approx 1.75$ on all duals.
- Deriving theorem-compliant step sizes from a weighted power iteration is a clean direction for the journal extension.

Backup: full-resolution RMSE comparison

Image RMSE under the paper's empirical step-size setup (paper-config $c_{hi} = 4$, $c_{lo} = 8$, $\sigma_{lo} = 4\sigma_{hi}$):

Resolution	iter 200			iter 500		
	Single	Two	Δ	Single	Two	Δ
512×512	0.1953	0.1628	-16.7%	0.1625	0.1317	-19.0%
256×256	0.1829	0.1192	-34.8%	0.0984	0.0768	-22.0%
128×128	0.0827	0.0515	-37.8%	0.0332	0.0128	-61.5%

- Iter-500 RMSE reduction of 19/22/61% at $512^2/256^2/128^2$.
- Improvement is largest at 128^2 ; the iter-200 column shows the gap is broad across image sizes during the “clinical-iteration” window.
- Caveat: image RMSE is still a weak summary at this angular range – the visual ladder is the right primary reading.

References

- E. Y. Sidky et al., *Accurate volume image reconstruction for DBT with directional-gradient and pixel sparsity regularization*, J. Med. Imaging 12(S1):S13013, 2025.
- A. Chambolle, T. Pock, *A first-order primal-dual algorithm for convex problems with applications to imaging*, J. Math. Imag. Vision 40(1):120–145, 2011.
- B. He, X. Yuan, *Convergence analysis of primal-dual algorithms for a saddle-point problem: from contraction perspective*, SIAM J. Imag. Sci. 5(1):119–149, 2012.
- Z. Zhang, B. Chen, D. Xia, E. Y. Sidky, X. Pan, *Directional-TV algorithm for image reconstruction from limited-angular-range data*, Med. Image Anal. 70:102030, 2021.
- I. Reiser, R. M. Nishikawa, *Task-based assessment of breast tomosynthesis*, Med. Phys. 37(4):1591–1600, 2010.